

Stochastic optimization algorithms

Lecture 4, 20200908

Evolutionary algorithms:
Background and introduction



Today's learning goals

- After this lecture you should be able to
 - Motivate the use of evolutionary algorithms.
 - Describe the central concepts in Darwinian evolution.
 - Describe the structure and functionality of genes (the processes of transcription and translation).
 - Summarize the main aspects of evolution.
 - Describe the basic functionality of evolutionary algorithms.
 - Exemplify some of the simplifications used in evolutionary algorithms (relative to biological evolution).


Why evolutionary algorithms?

- Evolutionary algorithms ...
 - ...do not (usually) get stuck in local optima.
 - ...can find several different and equally viable solutions to a given problem.
 - ... are easily applicable to a wide variety of problems.
 - ... but also require some experience for successful usage.


Example: Evolution of eyes




Demonstrations




Demonstrations




Demonstrations




Demonstrations




Demonstrations




Demonstrations



Demonstrations



Demonstrations



Today's learning goals

- After this lecture you should be able to
 - Motivate the use of evolutionary algorithms.
 - Describe the central concepts in Darwinian evolution.
 - Describe the structure and functionality of genes (the processes of transcription and translation).
 - Summarize the main aspects of evolution.
 - Describe the basic functionality of evolutionary algorithms.
 - Exemplify some of the simplifications used in evolutionary algorithms (relative to biological evolution).



Darwinian evolution

- Central concepts: Populations and individuals.
- A **population** is a group of **individuals** of the same **species**.
- Individuals of the same species can have fertile offspring.

Liger = tiger + lion!



Darwinian evolution

- Central concept: Gradual, hereditary change
 - **Gradualness**: Changes occur over very long times.
 - **Heredity**: Properties of an individual can be stored so that they can be transmitted to the next generation (reproduction).

Today's learning goals

- After this lecture you should be able to
 - Motivate the use of evolutionary algorithms.
 - Describe the central concepts in Darwinian evolution.
 - Describe the structure and functionality of genes (the processes of transcription and translation).
 - Summarize the main aspects of evolution.
 - Describe the basic functionality of evolutionary algorithms.
 - Exemplify some of the simplifications used in evolutionary algorithms (relative to biological evolution).





Darwinian evolution

- How is the transferable information stored?
 - ...in the **genome**, which (in humans) consist of 23 pairs of **chromosomes**.
 - Each chromosome, in turn, contains a large number of **genes**.

Chromosome numbers


- Fruit fly : 8
- Cat: 38
- Human: 46
- Ape: 48
- Horse: 64
- Dog: 78
- Carp: 104



Number of genes

- Approximate number of genes in selected species:
 - Bacterium: 500 – 6000
 - Yeast: 6000
 - Fruit fly: 13600
 - Human: 25000
- However, there are many species (e.g. some fish) with more genes than humans.

Chromosomes



Chromosome

Sequences of base pairs:

...A-C-C-T-C-G-G-T-C-A-G...

...T-G-G-A-G-C-C-A-G-T-C...

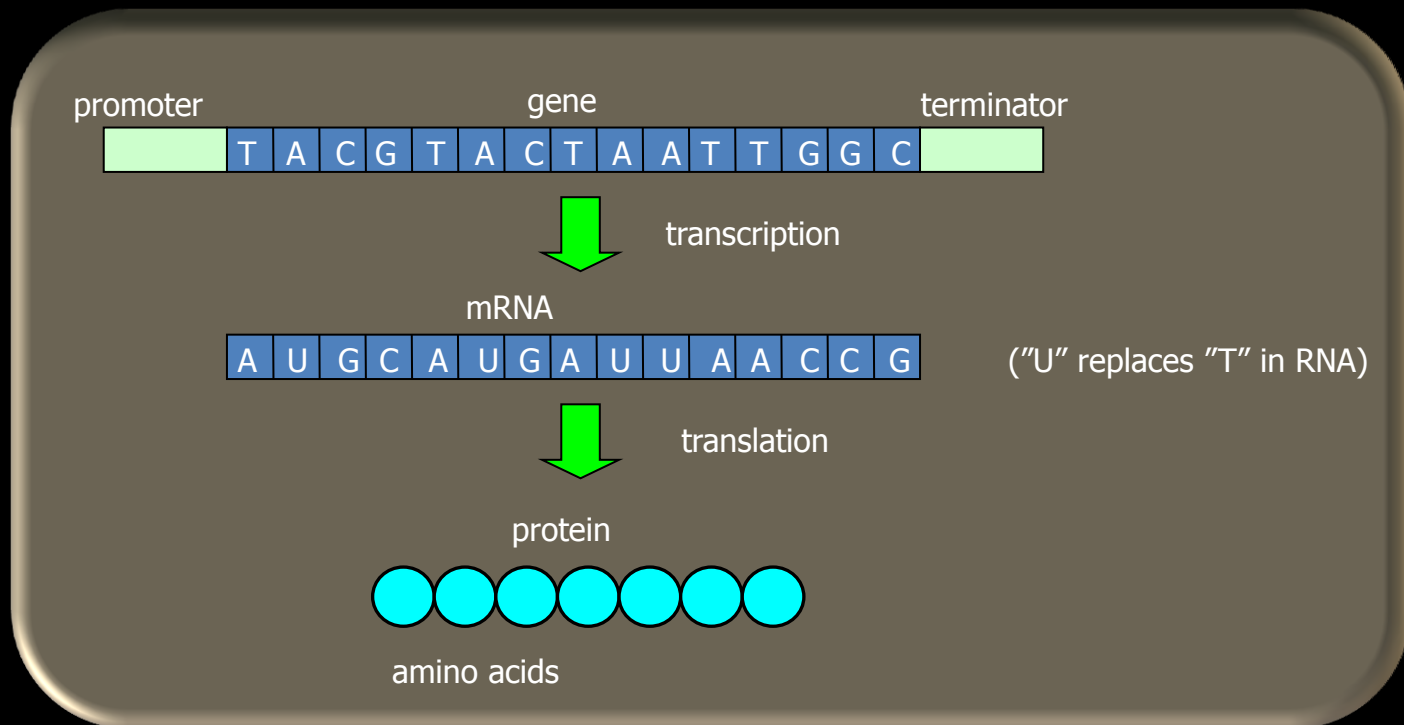
Genes

- Genes consist of sequences of **base pairs**, i.e. pairs of molecules (whose names are abbreviated A, C, T, G).
- A is always paired with T, and C is always paired with G:

... A C T G G T C A A T G ...
... T G A C C A G T T A C ...

The functionality of genes


- Genes are used for making proteins, through the two steps of transcription and translation:



pp. 35-39

Transcription


- DNA read by RNA polymerase => mRNA



pp. 35-39


Transcription

- U replaces T in mRNA



Translation

- mRNA used as a template to generate chains of amino acids (proteins):




pp. 35-39

Translation

- How does the system know *which* amino acid to pick?
 - 3-letter sequences code for an amino acid.
 - There are 20 (standard) amino acids => 3 letters required to encode all in base 4 (A, C, G, U).
 - Example: The sequence CUA codes for Leucine.
 - Some sequences (e.g. AUG) code for *start* or *stop*.

Translation



		Second Position					
		U	C	A	G		
First Position	U	UUU } Phe UUC } UUA } Leu UUG }	UCU } UCC } Ser UCA } UCG }	UAU } Tyr UAC } UAA Stop UAG Stop	UGU } Cys UGC } UGA Stop UGG Trp	U C A G	Third Position
	C	CUU } CUC } Leu CUA } CUG }	CCU } CCC } Pro CCA } CCG }	CAU } His CAC } CAA } Gln CAG }	CGU } CGC } Arg CGA } CGG }	U C A G	
	A	AUU } AUC } Ile AUA } AUG Met	ACU } ACC } Thr ACA } ACG }	AAU } Asn AAC } AAA } Lys AAG }	AGU } Ser AGC } AGA } Arg AGG }	U C A G	
	G	GUU } GUC } Val GUA } GUG }	GCU } GCC } Ala GCA } GCG }	GAU } Asp GAC } GAA } Glu GAG }	GGU } GGC } Gly GGA } GGG }	U C A G	

Today's learning goals

- After this lecture you should be able to
 - Motivate the use of evolutionary algorithms. ✓
 - Describe the central concepts in Darwinian evolution. ✓
 - Describe the structure and functionality of genes (the processes of transcription and translation). ✓
 - Summarize the main aspects of evolution.
 - Describe the basic functionality of evolutionary algorithms.
 - Exemplify some of the simplifications used in evolutionary algorithms (relative to biological evolution).

Additional concepts


- **Alleles:** The various settings (variants) of a gene
- **Genotype:** The genome with all its alleles
- **Phenotype:** The individual with all its traits, resulting from a given genotype.

Fitness, selection, and reproduction

- The most well-adapted individuals (i.e. those with highest **fitness**) are more likely to have the opportunity (**selection**) to **reproduce** (i.e. to spread their genetic material).
- Thus, even though selection is stochastic (i.e. contains a random element) it does not occur with equal probability for individuals – instead it strongly favors fit individuals.

Reproduction

- Asexual reproduction (e.g. bacteria):



- Sexual reproduction: Combines the genetic material of two individuals.

Mutations

- Mutations are small, random changes (copying errors)
- Many mutations have no effect at all, others are lethal.
- In general, even though mutations often have negative immediate effects, they give evolution new material to work (or, rather, experiment) with.

Evolution (summary)

- Evolution acts on populations of individuals (of a given species).
- Information is stored in the form of chromosomes, each consisting of many genes.
- Genes store the information needed to generate proteins (by transcription and translation).
- Well-adapted (fit) individuals spread their genetic material through reproduction.
- Mutations provide evolution with new opportunities.

pp. 35-39



Today's learning goals

- After this lecture you should be able to
 - Motivate the use of evolutionary algorithms. ✓
 - Describe the central concepts in Darwinian evolution. ✓
 - Describe the structure and functionality of genes (the processes of transcription and translation). ✓
 - Summarize the main aspects of evolution. ✓
 - Describe the basic functionality of evolutionary algorithms.
 - Exemplify some of the simplifications used in evolutionary algorithms (relative to biological evolution).

Evolutionary algorithms


- EA = evolutionary algorithm
- GA = genetic algorithm
- We will begin by describing GAs, which constitute an important special case.
- Algorithm 3.1 in the book.

pp. 40-45



Encoding

- Encoding: the variables of the problem are encoded in strings of digits known as chromosomes.



- The exact encoding procedure varies from problem to problem. Two examples (among many!):
 - Function optimization: Encode the variables x_1 , x_2 , x_3 etc.
 - Neural network optimization: Encode the network weights.


Step 1: Initialization

- Generate N random chromosomes (first **generation**):

1	0 1 0 1 1 0 1 1 1 0 1 0 1 1 1 0 1 1 0 0 0 1 0 1 0 1 1 0 1 0 0 1 1 0 1 0 1 1 1 1 0 0 0 1 0 1 0 0 1 0 1 1 0 1
2	1 0 1 1 1 0 1 0 0 0 1 0 1 1 0 0 0 1 0 0 0 0 0 1 1 0 0 0 1 1 0 1 0 1 1 0 0 1 0 1 0 0 0 0 1 0 1 1 0 0 0 1 1 0
3	0 1 0 1 1 1 0 1 1 1 1 1 0 0 1 0 1 1 0 0 0 0 1 0 0 1 1 0 1 0 0 1 1 0 1 0 1 0 0 0 1 1 1 0 1 0 1 1 0 0 0 1 0 1
	...
N	0 1 0 1 1 0 1 1 1 0 1 0 0 0 1 1 1 0 1 1 1 0 1 1 0 1 0 0 1 0 0 1 1 0 1 0 1 1 1 1 0 0 0 1 1 1 0 0 0 0 0 0 0 1

Step 2: Evaluation

- Step 2.1: Decode the chromosomes to obtain the variables.




- See example 3.1 (that illustrates the decoding of one chromosome).

Step 2.1: Decoding

- The first 10 bits are used to form the first variable.
- First, a value in the range $[0,1[$ is generated:

- $x_1^{\text{tmp}} = \sum_{j=1}^{10} 2^{-j} g_j$




- This value is then rescaled to the required range, which in this case happened to be $[-3,3]$:

- $x_1 = -3 + \frac{2 * 3}{1 - 2^{-10}} x_1^{\text{tmp}}$

Step 2: Decoding

- Next, x_2 is obtained in the same way, using the last 10 bits:



- $x_2^{\text{tmp}} = \sum_{j=1}^{10} 2^{-j} g_{j+10}$
- $x_2 = -3 + \frac{2 * 3}{1 - 2^{-10}} x_2^{\text{tmp}}$
- Thus, from the 20-bit chromosome, we have formed two variables, x_1 and x_2 , both in the range $[-3,3]$.

Step 2: Evaluation


- Step 2.2 Evaluate the individual, and assign a fitness score.
- Requirement: Fitness measure. Examples:
 - Function maximization: Use simply $f(x)$.
 - Function minimization: Use (e.g.) $1/f(x)$.
- See example 3.2.
- Step 2.3: Repeat Steps 2.1 and 2.2 for all individuals => a set of fitness scores is obtained, one for each individual.

Step 3: Form the next generation

- Step 3.1: Select two individuals (one at a time), such that individuals with high fitness have larger probability of selection.
- Various methods exist, e.g. Tournament selection, Roulette-wheel selection etc.
- (These methods will be studied in detail in the next lecture!)
- Selection is carried out with replacement, i.e. individuals can be selected more than once.

Example: Tournament selection


- Procedure for selection of one individual:
 1. Pick two individuals randomly from the population (equal probability for all individuals in the population)
 2. With probability p_{tour} (usually around 0.7-0.8) pick the better of the two individuals (higher fitness). Otherwise, pick the worse of the two.



Step 3: Form the next generation

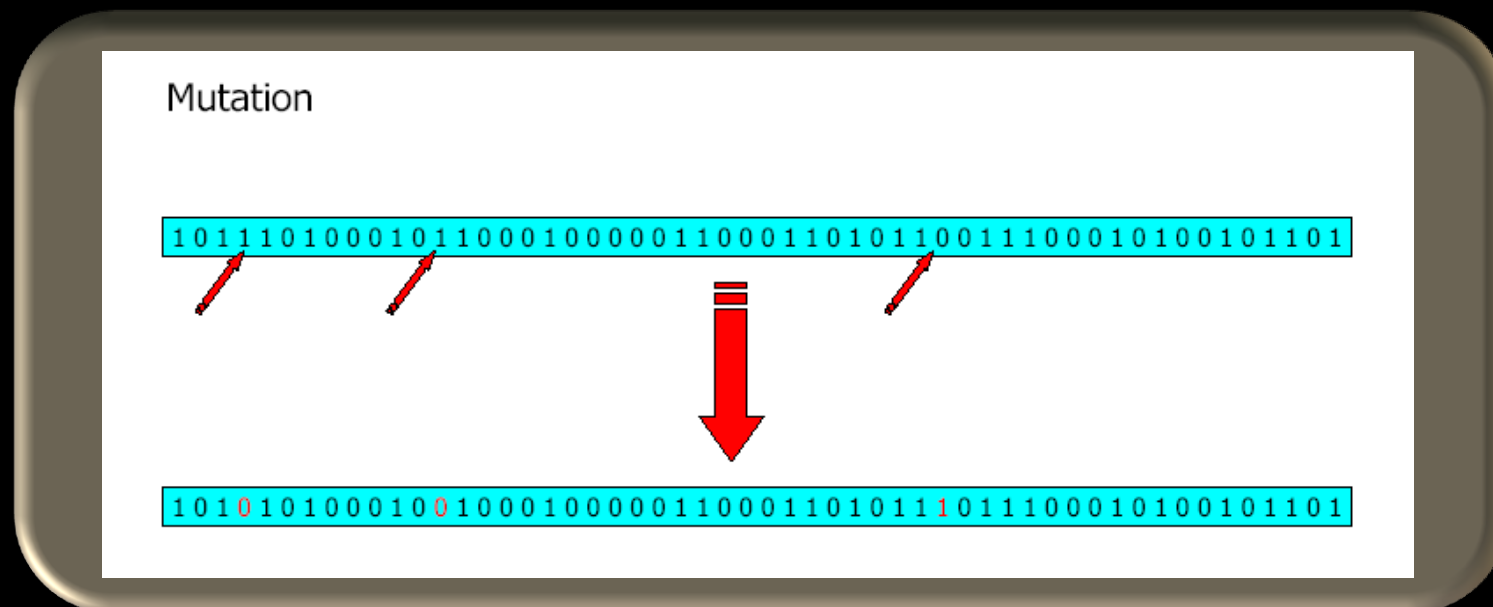
- Step 3.2: Generate two new chromosomes by crossing the chromosomes of the two selected individuals:

Crossover



Step 3: Form the next generation

- Step 3.3: Mutate the two new chromosomes:

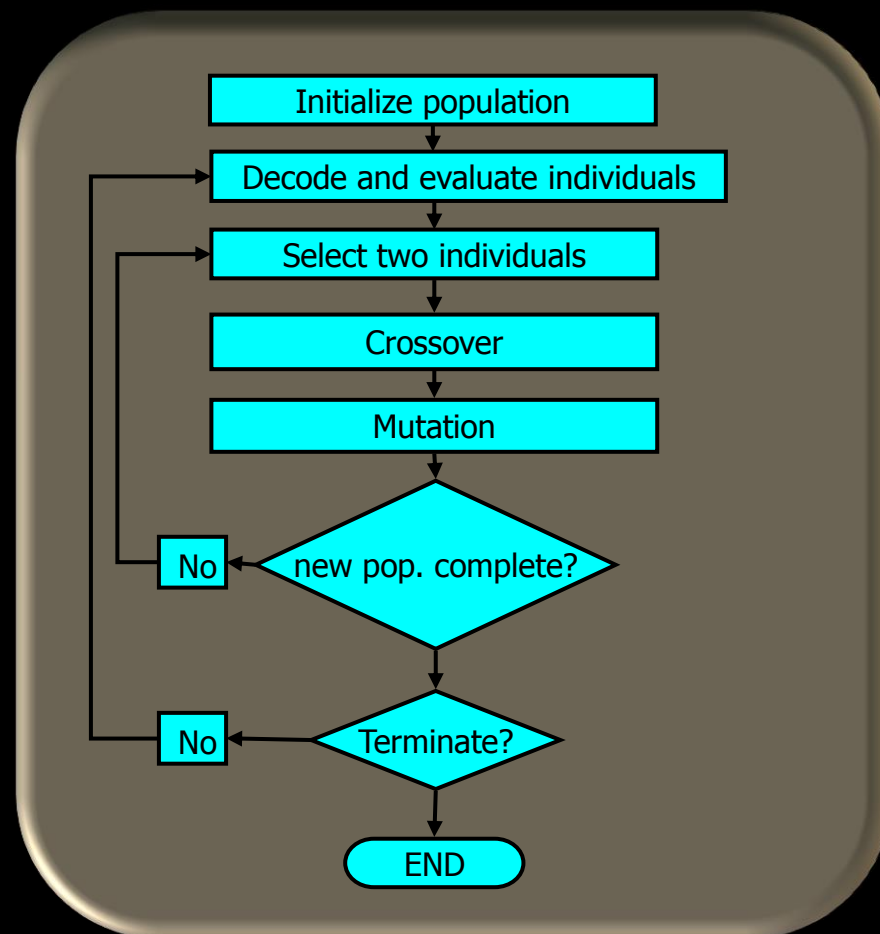


Step 3: Form the next generation

- Step 3.4: Repeat the entire procedure (Steps 3.1 – 3.3) $N/2$ times, to generate N new individuals. That is, do $N/2$ repetitions of...
 - ...selection of two individuals (for example by doing tournament selection twice),
 - ...crossover, and
 - ...mutation of the resulting pair of chromosomes
- Then replace the old set of chromosomes by the new ones.

Step 4: Repeat!

- The process is repeated (from Step 2) until a satisfactory solution has been found.



Example (Fig. 3.7)




Figure 3.7: The progress of a GA searching for the maximum of the function $f(x_1, x_2) = e^{-x_1^2 - x_2^2}$. The upper right panel shows the initial population, whereas the lower left and lower right panels show the population after two and 25 generations, respectively.

Today's learning goals

- After this lecture you should be able to
 - Motivate the use of evolutionary algorithms. ✓
 - Describe the central concepts in Darwinian evolution. ✓
 - Describe the structure and functionality of genes (the processes of transcription and translation). ✓
 - Summarize the main aspects of evolution. ✓
 - Describe the basic functionality of evolutionary algorithms. ✓
 - Exemplify some of the simplifications used in evolutionary algorithms (relative to biological evolution).


Evolution vs. evolutionary algorithms

- Evolutionary algorithms (EAs) are strongly simplified compared to biological evolution!
 - In EAs, chromosomes are used as lookup tables, i.e. there are no gene interactions etc.
 - In EAs, multicellularity (difference between cells etc.) is rarely considered. In biological systems, cells share the same genetic material, but develop differently (skin, bone, muscle etc.)
 - In EAs, there is no gene regulation. In biological systems, (many) genes are active, to varying degrees, during the lifetime of the individual.


Gene regulation

- Regulatory genes: (transcription factors) genes that regulate the expression of other genes. Example:

Repressor protein (= the product of some other (regulatory) gene) bound to operator site: transcription is prevented



Repressor *not* bound to operator: the RNA polymerase can reach the promoter and proceed with transcription:



Today's learning goals

- After this lecture you should be able to
 - Motivate the use of evolutionary algorithms. ✓
 - Describe the central concepts in Darwinian evolution. ✓
 - Describe the structure and functionality of genes (the processes of transcription and translation). ✓
 - Summarize the main aspects of evolution. ✓
 - Describe the basic functionality of evolutionary algorithms. ✓
 - Exemplify some of the simplifications used in evolutionary algorithms (relative to biological evolution). ✓

Home problem 1

- Consists of three parts.
 - Problems 1.1 and 1.3 are mandatory,
 - Problem 1.2 is voluntary (but needed for higher grades).
- Use Matlab when solving the home problems.
- Maximum total score: 10p.
- Should be solved (and handed in) *independently by each student*.
- Strict deadline: 20200922, 23.59.59.

Home problem 1

- Note: Read the instructions carefully (regarding the report, coding standard etc.)
- Before submitting:
 - What, **exactly**, is being asked in the problem(s)? Have you **answered** those questions?
 - Have you written and included a report in the required format?
 - Have you combined your programs *and* your report into ONE compressed file, in an allowed format (.zip or .7z)?
- Read the checklist (available on the web page, under the module called Miscellaneous).

Problem 1.1

- The penalty method:

In this problem, we shall use the penalty method (see pp. 30-33 in the course book) to find the minimum of the function


$$f(x_1, x_2) = (x_1 - 1)^2 + 2(x_2 - 2)^2, \quad (1)$$

subject to the constraint

$$g(x_1, x_2) = x_1^2 + x_2^2 - 1 \leq 0. \quad (2)$$

Problem 1.2

- (Classical) optimization under constraints



Problem 1.3

- Function optimization using genetic algorithms.
- Write a (*well-structured*) program for finding the minimum of the function:

$$g(x_1, x_2) = \left(1 + (x_1 + x_2 + 1)^2(19 - 14x_1 + 3x_1^2 - 14x_2 + 6x_1x_2 + 3x_2^2)\right) \times \\ \left(30 + (2x_1 - 3x_2)^2(18 - 32x_1 + 12x_1^2 + 48x_2 - 36x_1x_2 + 27x_2^2)\right)$$

Matlab introduction for SOAs

- This afternoon and evening. The document will be posted at 13.00, I will be online (Zoom) between 18.00 and 20.00.
- Important to attend this session!
- You will get a document with Matlab code describing the implementation of a genetic algorithm (GA).
- The GA will be applied in a function optimization task.