

# Stochastic optimization algorithms

## Lecture 9, 20200918

Neural networks and data analysis

# Note!



- The deadline for the home problem (HP1) is at 23.59.59 on Tuesday the 22<sup>nd</sup>.
- Read carefully the FAQ and the *checklist for home problem submission* (on the course web page) before submitting your solutions.
- Check also the comments for the IPP, and make any required adjustments.
- Penalties for delays, see the checklist.
- Next week, there is no lecture on Wednesday (23<sup>rd</sup>).

# Today's learning goals

- After this lecture you should be able to
  - Describe the biological background of neural networks.
  - Describe basic learning: Habituation and sensitization.
  - Compute the output of a feedforward neural network.
  - Apply a GA to optimize an artificial neural network.
  - Describe and discuss the concept of overfitting, as well as methods for avoiding it.


# Biological neurons

- The brains of animals consists of many (billions, in the case of humans and other mammals) **neurons**.




# Biological neurons

- Each **neuron** is connected to many (typically thousands) of other neurons via (mostly chemical) **synapses**.



# Biological neurons

- The signal flow within a neuron is electrical. The neuron makes a binary decision whether or not to fire a **spike**.



- Synapses can be either **excitatory** or **inhibitory**.

# Biological neurons

- There is a refractory period limiting the firing frequency to around 1 kHz.

---
- How can brains then carry out such complex tasks?
  - Parallel computation!
  - Humans:  $\sim 10^{12}$  neurons,  $\sim 10^{15}$  synapses.
  - $10^{12}$  neurons  $\times 10^3$  Hz  $\Rightarrow 10^{15}$  operations/s.
- As we shall see, the same principle, i.e. using many simple, interconnected computational units, is used in artificial neural networks.

# Biological neurons

- There is a refractory period limiting the firing frequency to around 1 kHz.
- How can brains then carry out such complex tasks?

---

  - Parallel computation!
  - Humans:  $\sim 10^{12}$  neurons,  $\sim 10^{15}$  synapses.
  - $10^{12}$  neurons  $\times 10^3$  Hz  $\Rightarrow 10^{15}$  operations/s.
- As we shall see, the same principle, i.e. using many simple, interconnected computational units, is used in artificial neural networks.



# Biological neurons

- There is a refractory period limiting the firing frequency to around 1 kHz.
- How can brains then carry out such complex tasks?
  - Parallel computation!
  - Humans:  $\sim 10^{12}$  neurons,  $\sim 10^{15}$  synapses.
  - $10^{12}$  neurons  $\times 10^3$  Hz  $\Rightarrow 10^{15}$  operations/s.
- As we shall see, the same principle, i.e. using many simple, interconnected computational units, is used in artificial neural networks.

# Biological neurons

- There is a refractory period limiting the firing frequency to around 1 kHz.
- How can brains then carry out such complex tasks?
  - Parallel computation!
  - Humans:  $\sim 10^{12}$  neurons,  $\sim 10^{15}$  synapses.
  - $10^{12}$  neurons  $\times 10^3$  Hz  $\Rightarrow 10^{15}$  operations/s.
- As we shall see, the same principle, i.e. using many simple, interconnected computational units, is used in artificial neural networks.

# Biological neurons

- There is a refractory period limiting the firing frequency to around 1 kHz.
- How can brains then carry out such complex tasks?
  - Parallel computation!
  - Humans:  $\sim 10^{12}$  neurons,  $\sim 10^{15}$  synapses.
  - $10^{12}$  neurons  $\times 10^3$  Hz  $\Rightarrow 10^{15}$  operations/s.
- As we shall see, the same principle, i.e. using many simple, interconnected computational units, is used in artificial neural networks.

# Biological neurons

- There is a refractory period limiting the firing frequency to around 1 kHz.
  - How can brains then carry out such complex tasks?
    - Parallel computation!
    - Humans:  $\sim 10^{12}$  neurons,  $\sim 10^{15}$  synapses.
    - $10^{12}$  neurons  $\times 10^3$  Hz  $\Rightarrow 10^{15}$  operations/s.
  - As we shall see, the same principle, i.e. using many simple, interconnected computational units, is used in artificial neural networks.
-

# Today's learning goals

- After this lecture you should be able to
  - Describe the biological background of neural networks.
  - Describe basic learning: Habituation and sensitization.
  - Compute the output of a feedforward neural network.
  - Apply a GA to optimize an artificial neural network.
  - Describe and discuss the concept of overfitting, as well as methods for avoiding it.



# Learning

- The information in neural networks is (mostly) stored in the connections (synapses) between neurons.
- Learning = modification of synapse strength.

# Learning

- The information in neural networks is (mostly) stored in the connections (synapses) between neurons.
  - Learning = modification of synapse strength.
-

# Learning

- The information in neural networks is (mostly) stored in the connections (synapses) between neurons.
- Learning = modification of synapse strength.
- Hebbian learning:

$$\frac{dw_{ij}}{dt} = \eta x_i x_j$$



# Learning

- The information in neural networks is (mostly) stored in the connections (synapses) between neurons.
- Learning = modification of synapse strength.
- Hebbian learning:

$$\frac{dw_{ij}}{dt} = \eta x_i x_j$$

- Modified hebbian learning:

$$\frac{dw_{ij}}{dt} = \eta (x_i - \bar{x}_i)(x_j - \bar{x}_j)$$

# Habituation and sensitization

- Habituation: Decreased response to neutral stimuli.
- Sensitization: Increased response to neutral stimuli following an aversive stimulus.
- Studied extensively by Kandel *et al.* (Nobel Prize 2000)
- Kandel studied the sea slug *Aplysia*, due to
  - ...its very thick neurons (easier to measure) and,
  - ...its rather small number of neurons (around 20000).

# Habituation and sensitization

- Habituation: Decreased response to neutral stimuli.
- Sensitization: Increased response to neutral stimuli following an aversive stimulus.
- Studied extensively by Kandel *et al.* (Nobel Prize 2000)
- Kandel studied the sea slug *Aplysia*, due to
  - ...its very thick neurons (easier to measure) and,
  - ...its rather small number of neurons (around 20000).

# Habituation and sensitization

- Habituation: Decreased response to neutral stimuli.
  - Sensitization: Increased response to neutral stimuli following an aversive stimulus.
  - Studied extensively by Kandel *et al.* (Nobel Prize 2000)
  - Kandel studied the sea slug *Aplysia*, due to
    - ...its very thick neurons (easier to measure) and,
    - ...its rather small number of neurons (around 20000).
-

# Habituation and sensitization



- Eric Kandel's Nobel lecture (highly recommended!):
  - <http://nobelprize.org/mediaplayer/?id=898>

# Habituation and sensitization

- Basic conclusions from Kandel's work
  - Short-term learning depends on changes in the release of neurotransmitters.
  - Long-term learning depends on changes in gene regulation that, in turn, modify synapses (more or less) permanently.

# Habituation and sensitization

- Basic conclusions from Kandel's work
    - Short-term learning depends on changes in the release of neurotransmitters.
    - Long-term learning depends on changes in gene regulation that, in turn, modify synapses (more or less) permanently.
-

# Today's learning goals


- After this lecture you should be able to
  - Describe the biological background of neural networks.
  - Describe basic learning: Habituation and sensitization.
  - Compute the output of a feedforward neural network.
  - Apply a GA to optimize an artificial neural network.
  - Describe and discuss the concept of overfitting, as well as methods for avoiding it.



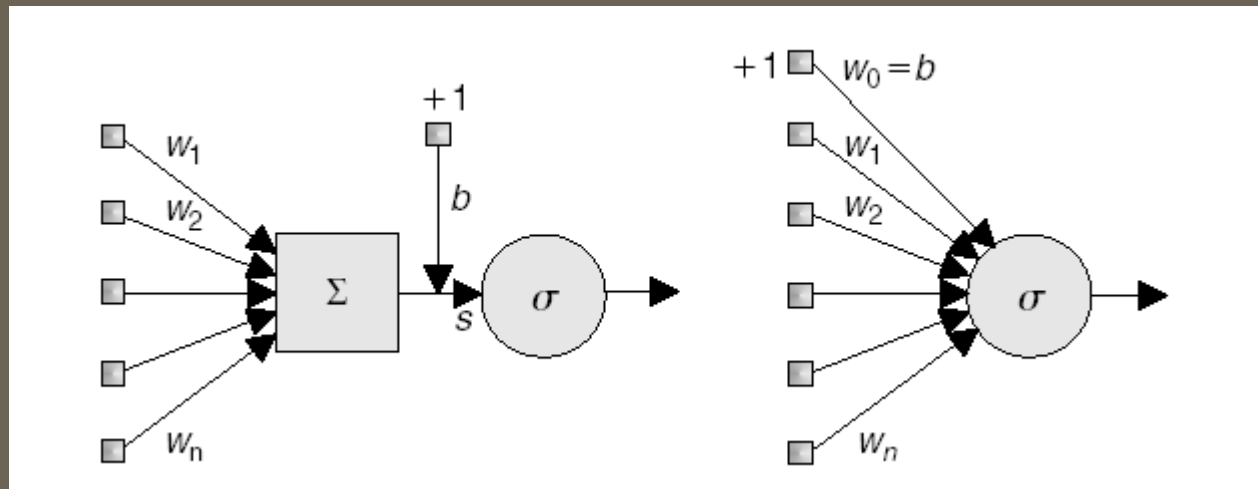


# Artificial neural networks

- Artificial neural networks (ANNs) consist of simple, interconnected computational units.
- Two kinds:
  - Feedforward neural networks (FFNNs)
  - Recurrent neural networks (RNNs)

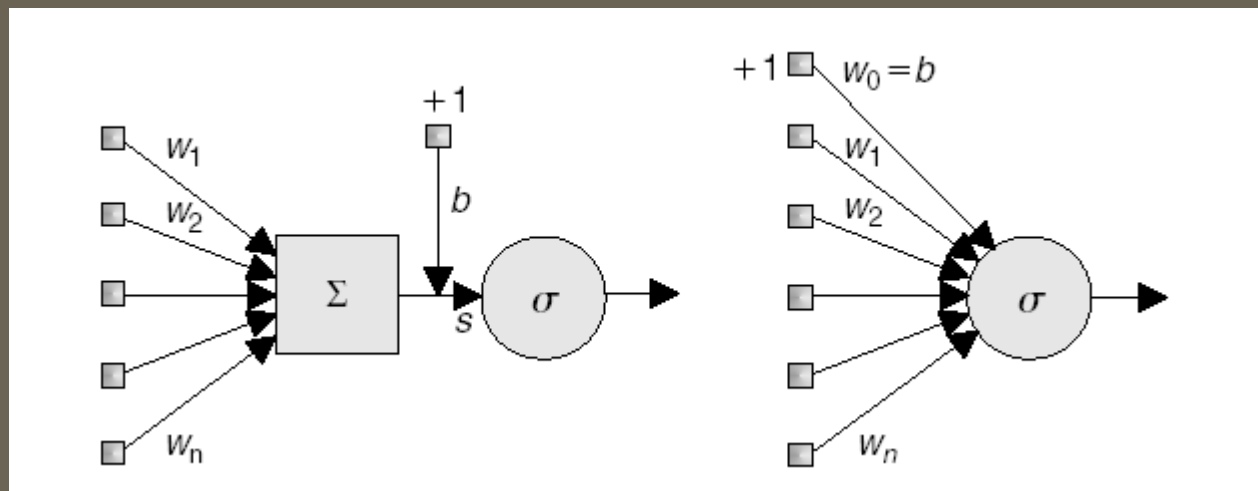


# Neurons (in ANNs)



- Output:  $y = \sigma\left(\sum_{j=1}^n w_j x_j + b\right) \equiv \sigma\left(\sum_{j=0}^n w_j x_j\right)$


# Neurons (in ANNs)




- Output:  $y = \sigma\left(\sum_{j=1}^n w_j x_j + b\right) \equiv \sigma\left(\sum_{j=0}^n w_j x_j\right)$

For example the logistic sigmoid  $\sigma(z) = 1/(1 + e^{-cz})$ , where  $c$  is a constant.

# Artificial neural networks




# Example A.1




# Example A.1

Input unit (simply  
mediates signal –  
no other computation).



# Example A.1

Input unit (simply  
mediates signal –  
no other computation).




Neuron – computation as  
in the previous slide.

# Example A.1

Input unit (simply mediates signal – no other computation).


Neuron – computation as in the previous slides.



- You should know how to compute the output of an FFNN, but you do not need to know backpropagation (for this course, at least).




# Example A.1



- You should know how to compute the output of an FFNN, but you do not need to know backpropagation (for this course, at least).

# Example A.1




Output from the first  
neuron in the hidden layer:

$$y_1^H = \sigma \left( \sum_{p=0}^2 w_{1p}^{I \rightarrow H} y_p^I \right) = \sigma(0.2 \times 1.0 + 0.3 \times 1.0 - 0.1 \times 0)$$

$$= \sigma(0.5) = \frac{1}{1 + e^{-0.5}} = 0.6225,$$

# Example A.1




Output from the first neuron in the hidden layer:

$$y_1^H = \sigma \left( \sum_{p=0}^2 w_{1p}^{I \rightarrow H} y_p^I \right) = \sigma(0.2 \times 1.0 + 0.3 \times 1.0 - 0.1 \times 0)$$

$$= \sigma(0.5) = \frac{1}{1 + e^{-0.5}} = 0.6225,$$

# Example A.1




Output from the first neuron in the hidden layer:

$$y_1^H = \sigma \left( \sum_{p=0}^2 w_{1p}^{I \rightarrow H} y_p^I \right) = \sigma(0.2 \times 1.0 + 0.3 \times 1.0 - 0.1 \times 0)$$

$$= \sigma(0.5) = \frac{1}{1 + e^{-0.5}} = 0.6225,$$

# Example A.1




Output from the first neuron in the hidden layer:

$$y_1^H = \sigma \left( \sum_{p=0}^2 w_{1p}^{I \rightarrow H} y_p^I \right) = \sigma(0.2 \times 1.0 + 0.3 \times 1.0 - 0.1 \times 0)$$

$$= \sigma(0.5) = \frac{1}{1 + e^{-0.5}} = 0.6225,$$

# Example A.1




Output from the first neuron in the hidden layer:

$$y_1^H = \sigma \left( \sum_{p=0}^2 w_{1p}^{I \rightarrow H} y_p^I \right) = \sigma(0.2 \times 1.0 + 0.3 \times 1.0 - 0.1 \times 0)$$

$$= \sigma(0.5) = \frac{1}{1 + e^{-0.5}} = 0.6225,$$

In this example,  
 $\sigma(z) = 1/(1 + e^{-cz})$ ,  
 with  $c = 1$ .

# Example A.1




Output from the second neuron in the hidden layer:

$$y_2^H = \sigma \left( \sum_{p=0}^2 w_{2p}^{I \rightarrow H} y_p^I \right) = \sigma(-0.2 \times 1.0 - 0.2 \times 1.0 + 0.1 \times 0)$$

$$= \sigma(-0.4) = \frac{1}{1 + e^{0.4}} = 0.4013.$$

# Example A.1




Output from the neuron in the output layer:

$$y_1^O = \sigma \left( \sum_{s=0}^2 w_{1s}^{H \rightarrow O} y_s^H \right) = \sigma(0.2 \times 1.0 + 0.5 \times 0.6225 + 0.1 \times 0.4013) \\ = \sigma(0.5514) = 0.6345.$$



# Example A.1



Output from the neuron in the output layer:

$$y_1^O = \sigma \left( \sum_{s=0}^2 w_{1s}^{H \rightarrow O} y_s^H \right) = \sigma(0.2 \times 1.0 + 0.5 \times 0.6225 + 0.1 \times 0.4013)$$

$$= \sigma(0.5514) = 0.6345.$$

$y_1^H$        $y_2^H$

# Today's learning goals

- After this lecture you should be able to
  - Describe the biological background of neural networks.
  - Describe basic learning: Habituation and sensitization.
  - Compute the output of a feedforward neural network.
  - Apply a GA to optimize an artificial neural network.
  - Describe and discuss the concept of overfitting, as well as methods for avoiding it.



# Training methods


- **Strongly guided learning:** input-output pairs available, feedback (error signal) given (to the network) for every input-output pair.
- **Weakly guided learning:** No input-output pairs. Instead, the feedback is given at (for example) the end of a lengthy simulation (see, for instance, the robot example (videos) from Lecture 1).

# Using a GA to optimize an ANN


- In many cases, the data set consists of input-output pairs, such that an error signal can be formed, based on the difference between the desired and actual outputs.
- In those cases, one can apply backpropagation (which requires the gradient of the error signal).
- In other cases, one might not have a set of input-output pairs – Instead feedback (output) may only be given, say, at the end of a long evaluation => backpropagation is not suitable.
- In such cases, one can apply a GA to train the network.

# Using a GA to optimize an ANN

Example: 2-2-1 network (6 weights (dark blue), 3 biases (red))




Encode in chromosomes with 9 real-valued genes:




Evaluate and assign fitness values, then apply selection, crossover and mutation (i.e. as usual). The evaluation varies from case to case, of course.

# Using a GA to optimize an ANN

Example: 2-2-1 network (6 weights (dark blue), 3 biases (red))




Encode in chromosomes with 9 real-valued genes:



Evaluate and assign fitness values, then apply selection, crossover and mutation (i.e. as usual). The evaluation varies from case to case, of course.

# Using a GA to optimize an ANN

Example: 2-2-1 network (6 weights (dark blue), 3 biases (red))




Encode in chromosomes with 9 real-valued genes:




Evaluate and assign fitness values, then apply selection, crossover and mutation (i.e. as usual). The evaluation varies from case to case, of course.

# Using a GA to optimize an ANN

Example: 2-2-1 network (6 weights (dark blue), 3 biases (red))



Encode in chromosomes with 9 real-valued genes:



Evaluate and assign fitness values, then apply selection, crossover and mutation (i.e. as usual). The evaluation varies from case to case, of course.



# Today's learning goals

- After this lecture you should be able to
  - Describe the biological background of neural networks.
  - Describe basic learning: Habituation and sensitization.
  - Compute the output of a feedforward neural network.
  - Apply a GA to optimize an artificial neural network.
  - Describe and discuss the concept of overfitting, as well as methods for avoiding it.



# Experiment design and overfitting

- When fitting a function (or a neural network, or some other structure) to a data set of limited size, there is a risk of fitting the noise => decrease in predictive power.
- This (fitting noise) is called **overfitting**.
- To avoid overfitting, one typically divides the available data into three sets:
  - Training
  - Validation (for determining when to quit training)
  - Test (for final evaluation, not used during training)

# Experiment design and overfitting

- When fitting a function (or a neural network, or some other structure) to a data set of limited size, there is a risk of fitting the noise => decrease in predictive power.
- This (fitting noise) is called **overfitting**.
- To avoid overfitting, one typically divides the available data into three sets:
  - Training
  - Validation (for determining when to quit training)
  - Test (for final evaluation, not used during training)

# Experiment design and overfitting

- When fitting a function (or a neural network, or some other structure) to a data set of limited size, there is a risk of fitting the noise => decrease in predictive power.
- This (fitting noise) is called **overfitting**.
- To avoid overfitting, one typically divides the available data into three sets:

---

  - Training
  - Validation (for determining when to quit training)
  - Test (for final evaluation, not used during training)

# Experiment design and overfitting

- When fitting a function (or a neural network, or some other structure) to a data set of limited size, there is a risk of fitting the noise => decrease in predictive power.
- This (fitting noise) is called **overfitting**.
- To avoid overfitting, one typically divides the available data into three sets:
  - Training
  - Validation (for determining when to quit training)
  - Test (for final evaluation, not used during training)


# Experiment design and overfitting

- When fitting a function (or a neural network, or some other structure) to a data set of limited size, there is a risk of fitting the noise => decrease in predictive power.
- This (fitting noise) is called **overfitting**.
- To avoid overfitting, one typically divides the available data into three sets:
  - Training
  - Validation (for determining when to quit training)
  - Test (for final evaluation, not used during training)

# Experiment design and overfitting


- When fitting a function (or a neural network, or some other structure) to a data set of limited size, there is a risk of fitting the noise => decrease in predictive power.
- This (fitting noise) is called **overfitting**.
- To avoid overfitting, one typically divides the available data into three sets:
  - Training
  - Validation (for determining when to quit training)
  - Test (for final evaluation, not used during training)

# Experiment design and overfitting






# Experiment design and overfitting



# Experiment design and overfitting



# Procedure (avoiding overfitting)

- Use the **training data** to give feedback to the training algorithm.
- Use the **validation data** to determine when to stop the training – measure the error (or fitness) over the validation set, but do *not* make the result of the measurement available to the training algorithm.
- Use the (previously unused) **test data** (after completing the training) to test the results (on the selected system, i.e. the system with best validation performance).

# Procedure (avoiding overfitting)

- Use the **training data** to give feedback to the training algorithm.
- Use the **validation data** to determine when to stop the training – measure the error (or fitness) over the validation set, but do *not* make the result of the measurement available to the training algorithm.
- Use the (previously unused) **test data** (after completing the training) to test the results (on the selected system, i.e. the system with best validation performance).


# Procedure (avoiding overfitting)

- Use the **training data** to give feedback to the training algorithm.
- Use the **validation data** to determine when to stop the training – measure the error (or fitness) over the validation set, but do *not* make the result of the measurement available to the training algorithm.
- Use the (previously unused) **test data** (after completing the training) to test the results (on the selected system, i.e. the system with best validation performance).


# Example: Time series prediction

- Example C.4: Using a neural network to predict annual rainfall.
- In this case, each data point consists of an input-output pair:
  - Input: Rainfall in years  $T-1$ ,  $T-2$ ,  $T-3$ ,  $T-4$ ,  $T-5$
  - Output: Rainfall in year  $T$ .
- A total of 90 training data points and 31 validation data points.
- Objective: Find an FFNN whose output differs as little as possible from the actual output.

# Example: Time series prediction




# Example: Time series prediction






# Example: Time series prediction



# Example: Time series prediction

Best validation performance



# Today's learning goals

- After this lecture you should be able to
  - Describe the biological background of neural networks.
  - Describe basic learning: Habituation and sensitization.
  - Compute the output of a feedforward neural network.
  - Apply a GA to optimize an artificial neural network.
  - Describe and discuss the concept of overfitting, as well as methods for avoiding it.

